



Robotics


Path Planning

Path finding vs. trajectory optimization, local vs. global, Dijkstra, Probabilistic Roadmaps, Rapidly Exploring Random Trees, non-holonomic systems, car system equation, path-finding for non-holonomic systems, control-based sampling, Dubins curves



Marc Toussaint
University of Stuttgart
Winter 2016/17

Lecturer: Peter Englert


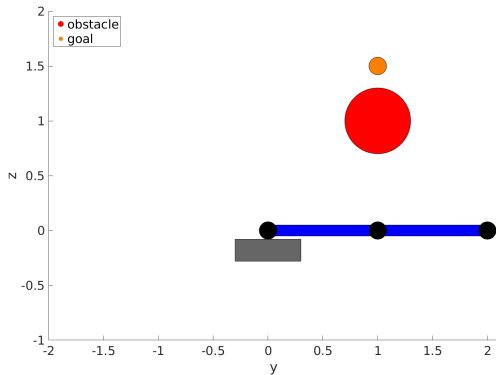
Path finding examples




Path finding examples



Path finding examples



Trajectory from initial state to goal:




Path finding examples




Alpha-Puzzle, solved with James Kuffner's RRTs

Path finding examples




J. Kuffner, K. Nishiwaki, S. Kagami, M. Inaba, and H. Inoue. Footstep Planning Among Obstacles for Biped Robots. Proc. IEEE/RSJ Int. Conf. on Intelligent Robots and Systems (IROS), 2001.

Path finding examples



S. M. LaValle and J. J. Kuffner. Randomized Kinodynamic Planning.
International Journal of Robotics Research, 20(5):378–400, May 2001.

Feedback control, path finding, trajectory optim.




- Feedback Control: E.g., $q_{t+1} = q_t + J^\sharp(y^* - \phi(q_t))$
- Trajectory Optimization: $\operatorname{argmin}_{q_{0:T}} f(q_{0:T})$
- Path Finding: Find some $q_{0:T}$ with only valid configurations

Outline

- **Really briefly:** Heuristics & Discretization (slides from Howie Choset's CMU lectures)
 - Bugs algorithm
 - Potentials to guide feedback control
 - Dijkstra
- **Sample-based Path Finding**
 - Probabilistic Roadmaps
 - Rapidly Exploring Random Trees
- **Non-holonomic systems**

Background

A better bug?




"Bug 2" Algorithm

- 1) head toward goal on the *m-line*
- 2) if an obstacle is in the way, follow it until you encounter the *m-line* again.
- 3) Leave the obstacle and continue toward the goal

A better bug?


"Bug 2" Algorithm



- 1) head toward goal on the *m-line*
- 2) if an obstacle is in the way, follow it until you encounter the *m-line* again.
- 3) Leave the obstacle and continue toward the goal

A better bug?

"Bug 2" Algorithm



- 1) head toward goal on the *m-line*
- 2) if an obstacle is in the way, follow it until you encounter the m-line again *closer to the goal*.
- 3) Leave the obstacle and continue toward the goal

Better or worse than Bug1?




BUG algorithms – conclusions


- Other variants: TangentBug, VisBug, RoverBug, WedgeBug, ...
- only 2D! (TangentBug has extension to 3D)
- Guaranteed convergence
- Still active research:
K. Taylor and S.M. LaValle: *I-Bug: An Intensity-Based Bug Algorithm*

⇒ Useful for minimalistic, robust 2D goal reaching
– not useful for finding paths in joint space

Start-Goal Algorithm: Potential Functions



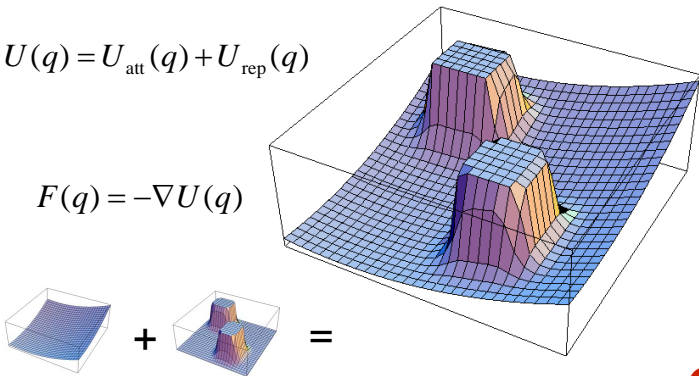
Repulsive Potential




Total Potential Function

$$U(q) = U_{\text{att}}(q) + U_{\text{rep}}(q)$$

$$F(q) = -\nabla U(q)$$



Local Minimum Problem with the Charge Analogy



Potential fields – conclusions

- Very simple, therefore popular
- In our framework: Combining a goal (endeffector) task variable, with a constraint (collision avoidance) task variable; then using inv. kinematics is *exactly* the same as “Potential Fields”

⇒ Does not solve locality problem of feedback control.

The Wavefront in Action (Part 1)

- Starting with the goal, set all adjacent cells with “0” to the current cell + 1
 - 4-Point Connectivity or 8-Point Connectivity?
 - Your Choice We'll use 8-Point Connectivity in our example

7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	1	1	1	1	1	1	1	1	0	0	0	0
3	0	0	0	0	1	1	1	1	1	1	1	1	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	3	3
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	3	2
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15



The Wavefront in Action (Part 2)

- Now repeat with the modified cells
 - This will be repeated until no 0's are adjacent to cells with values ≥ 2
 - 0's will only remain when regions are unreachable

7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	1	1	1	1	1	1	1	1	0	0	0	0
3	0	0	0	0	1	1	1	1	1	1	1	1	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	0	0	0	4	4	4
1	0	0	0	0	0	0	0	0	0	0	0	0	0	4	3	3
0	0	0	0	0	0	0	0	0	0	0	0	0	0	4	3	2
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15



The Wavefront in Action (Part 3)

- Repeat again...

7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	1	1	1	1	1	1	1	1	0	0	0
3	0	0	0	0	1	1	1	1	1	1	1	1	5	5	5
2	0	0	0	0	0	0	0	0	0	0	0	0	5	4	4
1	0	0	0	0	0	0	0	0	0	0	0	0	5	4	3
0	0	0	0	0	0	0	0	0	0	0	0	0	5	4	2
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14

The Wavefront in Action (Part 4)

- And again...

7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	1	1	1	1	1	1	1	1	6	6	6
3	0	0	0	0	1	1	1	1	1	1	1	1	5	5	5
2	0	0	0	0	0	0	0	0	0	0	0	6	5	4	4
1	0	0	0	0	0	0	0	0	0	0	0	6	5	4	3
0	0	0	0	0	0	0	0	0	0	0	0	6	5	4	2
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14

The Wavefront in Action (Part 5)

- And again until...

7	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	0	0	0	7	7	7	7	
4	0	0	0	0	1	1	1	1	1	1	1	1	6	6	6	
3	0	0	0	0	1	1	1	1	1	1	1	1	5	5	5	
2	0	0	0	0	0	0	0	0	0	0	7	6	5	4	4	
1	0	0	0	0	0	0	0	0	0	0	7	6	5	4	3	
0	0	0	0	0	0	0	0	0	0	0	7	6	5	4	2	
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

The Wavefront in Action (Done)


- You're done
 - Remember, 0's should only remain if unreachable regions exist

7	18	17	16	15	14	13	12	11	10	9	9	9	9	9	9	9
6	17	17	16	15	14	13	12	11	10	9	8	8	8	8	8	8
5	17	16	16	15	14	13	12	11	10	9	8	7	7	7	7	7
4	17	16	15	15	1	1	1	1	1	1	1	1	6	6	6	6
3	17	16	15	14	1	1	1	1	1	1	1	1	5	5	5	5
2	17	16	15	14	13	12	11	10	9	8	7	6	5	4	4	4
1	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	3
0	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

The Wavefront, Now What?

- To find the shortest path, according to your metric, simply always move toward a cell with a lower number
 - The numbers generated by the Wavefront planner are roughly proportional to their distance from the goal

Two
possible
shortest
paths
shown




Dijkstra Algorithm

- Is efficient in **discrete domains**
 - Given start and goal node in an arbitrary graph
 - Incrementally label nodes with their distance-from-start
- Produces optimal (shortest) paths
- Applying this to continuous domains requires discretization
 - Grid-like discretization in high-dimensions is daunting! (*curse of dimensionality*)
 - What are other ways to “discretize” space more efficiently?


Sample-based Path Finding

Probabilistic Road Maps



[Kavraki, Svetska, Latombe, Overmars, 95]

Probabilistic Road Maps




[Kavraki, Svetska, Latombe, Overmars, 95]

$q \in \mathbb{R}^n$ describes configuration

Q_{free} is the set of configurations without collision


Probabilistic Road Maps



[Kavraki, Svetska, Latombe, Overmars, 95]

- Probabilistic Road Maps generate a graph $G = (V, E)$ of configurations
 - such that configurations along each edges are $\in Q_{free}$

Probabilistic Road Maps



Given the graph, use (e.g.) Dijkstra to find path from q_{start} to q_{goal} .

Probabilistic Road Maps – generation


Input: number n of samples, number k number of nearest neighbors

Output: PRM $G = (V, E)$

```
1: initialize  $V = \emptyset, E = \emptyset$ 
2: while  $|V| < n$  do                                     // find  $n$  collision free points  $q_i$ 
3:    $q \leftarrow$  random sample from  $Q$ 
4:   if  $q \in Q_{\text{free}}$  then  $V \leftarrow V \cup \{q\}$ 
5: end while
6: for all  $q \in V$  do                                       // check if near points can be connected
7:    $N_q \leftarrow k$  nearest neighbors of  $q$  in  $V$ 
8:   for all  $q' \in N_q$  do
9:     if  $\text{path}(q, q') \in Q_{\text{free}}$  then  $E \leftarrow E \cup \{(q, q')\}$ 
10:  end for
11: end for
```


where $\text{path}(q, q')$ is a local planner (easiest: straight line)

Local Planner



tests collisions up to a specified resolution δ

Problem: Narrow Passages



The smaller the gap (clearance ϱ) the more unlikely to sample such points.

PRM theory

(for uniform sampling in d -dim space)

- Let $a, b \in Q_{\text{free}}$ and γ a path in Q_{free} connecting a and b .

Then the probability that *PRM* found the path after n samples is

$$P(\text{PRM-success} \mid n) \geq 1 - \frac{2|\gamma|}{\varrho} e^{-\sigma \varrho^d n}$$


$$\sigma = \frac{|B_1|}{2^d |Q_{\text{free}}|}$$

ϱ = clearance of γ (distance to obstacles)


(roughly: the exponential term are “volume ratios”)

- This result is called *probabilistic complete* (one can achieve any probability with high enough n)
- For a given success probability, n needs to be exponential in d


Other PRM sampling strategies



uniform



Gaussian




Bridge


illustration from O. Brock's lecture

Gaussian: $q_1 \sim \mathcal{U}$; $q_2 \sim \mathcal{N}(q_1, \sigma)$; if $q_1 \in Q_{\text{free}}$ and $q_2 \notin Q_{\text{free}}$, add q_1 (or vice versa).


Other PRM sampling strategies



uniform



Gaussian




Bridge

illustration from O. Brock's lecture


Gaussian: $q_1 \sim \mathcal{U}$; $q_2 \sim \mathcal{N}(q_1, \sigma)$; if $q_1 \in Q_{\text{free}}$ and $q_2 \notin Q_{\text{free}}$, add q_1 (or vice versa).

Bridge: $q_1 \sim \mathcal{U}$; $q_2 \sim \mathcal{N}(q_1, \sigma)$; $q_3 = (q_1 + q_2)/2$; if $q_1, q_2 \notin Q_{\text{free}}$ and $q_3 \in Q_{\text{free}}$, add q_3 .


Other PRM sampling strategies



uniform



Gaussian



Bridge

illustration from O. Brock's lecture

Gaussian: $q_1 \sim \mathcal{U}$; $q_2 \sim \mathcal{N}(q_1, \sigma)$; if $q_1 \in Q_{\text{free}}$ and $q_2 \notin Q_{\text{free}}$, add q_1 (or vice versa).

Bridge: $q_1 \sim \mathcal{U}$; $q_2 \sim \mathcal{N}(q_1, \sigma)$; $q_3 = (q_1 + q_2)/2$; if $q_1, q_2 \notin Q_{\text{free}}$ and $q_3 \in Q_{\text{free}}$, add q_3 .

- Sampling strategy can be made more intelligence: “utility-based sampling”
- Connection sampling (once earlier sampling has produced connected components)

Probabilistic Roadmaps – conclusions


- Pros:
 - Algorithmically very simple
 - Highly explorative
 - Allows probabilistic performance guarantees
 - Good to answer many queries in an *unchanged* environment
- Cons:
 - Precomputation of exhaustive roadmap takes a long time (but not necessary for “Lazy PRMs”)

Rapidly Exploring Random Trees


2 motivations:

- Single Query path finding: Focus computational efforts on paths for specific $(q_{\text{start}}, q_{\text{goal}})$
- Use actually controllable DoFs to incrementally explore the search space: *control-based* path finding.

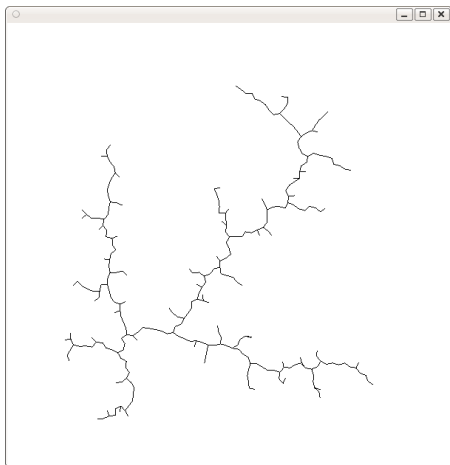
(Ensures that RRTs can be extended to handling differential constraints.)




$$n = 1$$




$n = 100$




$n = 300$



$n = 600$



$n = 1000$



$n = 2000$

Rapidly Exploring Random Trees

Simplest RRT with straight line local planner and step size α

Input: q_{start} , number n of nodes, stepsize α

Output: tree $T = (V, E)$

- 1: initialize $V = \{q_{\text{start}}\}$, $E = \emptyset$
 - 2: **for** $i = 0 : n$ **do**
 - 3: $q_{\text{target}} \leftarrow$ random sample from Q
 - 4: $q_{\text{near}} \leftarrow$ nearest neighbor of q_{target} in V
 - 5: $q_{\text{new}} \leftarrow q_{\text{near}} + \frac{\alpha}{|q_{\text{target}} - q_{\text{near}}|} (q_{\text{target}} - q_{\text{near}})$
 - 6: **if** $q_{\text{new}} \in Q_{\text{free}}$ **then** $V \leftarrow V \cup \{q_{\text{new}}\}$, $E \leftarrow E \cup \{(q_{\text{near}}, q_{\text{new}})\}$
 - 7: **end for**
-


Rapidly Exploring Random Trees

RRT growing directedly towards the goal


Input: q_{start} , q_{goal} , number n of nodes, stepsize α , β

Output: tree $T = (V, E)$

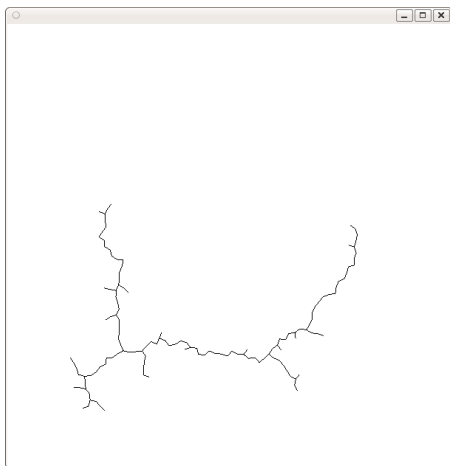
```
1: initialize  $V = \{q_{\text{start}}\}$ ,  $E = \emptyset$ 
2: for  $i = 0 : n$  do
3:   if  $\text{rand}(0, 1) < \beta$  then  $q_{\text{target}} \leftarrow q_{\text{goal}}$ 
4:   else  $q_{\text{target}} \leftarrow$  random sample from  $Q$ 
5:    $q_{\text{near}} \leftarrow$  nearest neighbor of  $q_{\text{target}}$  in  $V$ 
6:    $q_{\text{new}} \leftarrow q_{\text{near}} + \frac{\alpha}{|q_{\text{target}} - q_{\text{near}}|} (q_{\text{target}} - q_{\text{near}})$ 
7:   if  $q_{\text{new}} \in Q_{\text{free}}$  then  $V \leftarrow V \cup \{q_{\text{new}}\}$ ,  $E \leftarrow E \cup \{(q_{\text{near}}, q_{\text{new}})\}$ 
8: end for
```



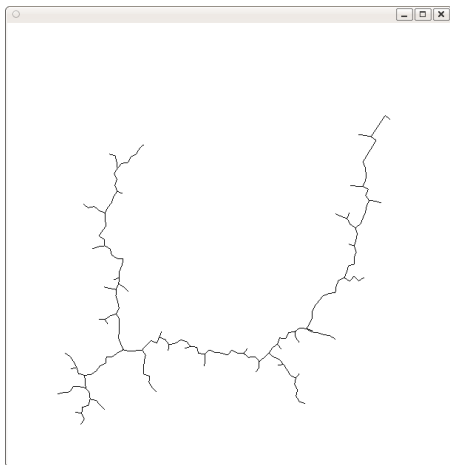
$$n = 1$$



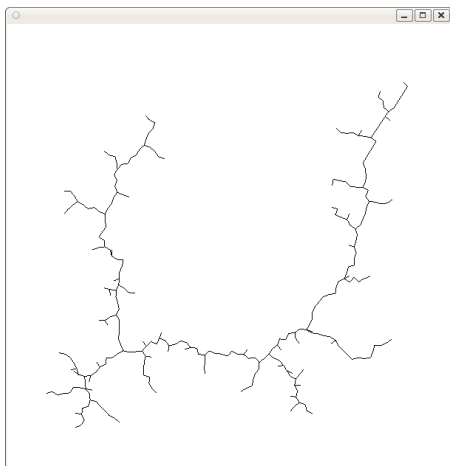
$$n = 100$$



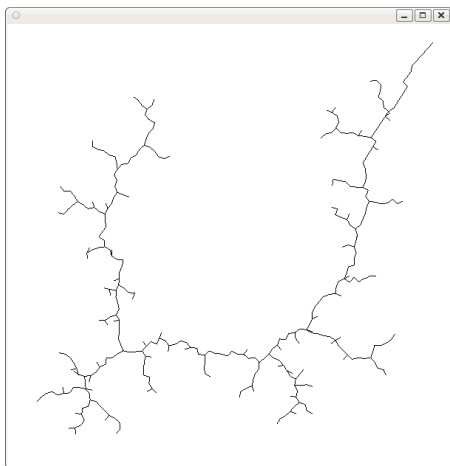
$$n = 200$$



$n = 300$



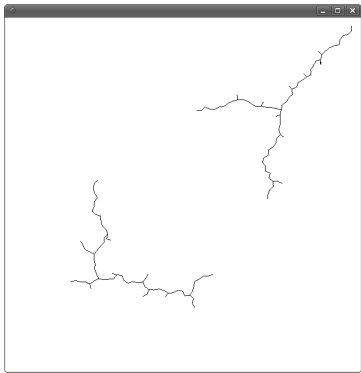
$n = 400$



$n = 500$

Bi-directional search

- grow two trees starting from q_{start} and q_{goal}



let one tree grow towards the other
(e.g., “choose q_{new} of T_1 as q_{target} of T_2 ”)

Summary: RRTs

- Pros (shared with PRMs):
 - Algorithmically very simple
 - Highly explorative
 - Allows probabilistic performance guarantees
- Pros (beyond PRMs):
 - Focus computation on single query ($q_{\text{start}}, q_{\text{goal}}$) problem
 - Trees from multiple queries can be merged to a roadmap
 - Can be extended to differential constraints (nonholonomic systems)
- To keep in mind (shared with PRMs):
 - The metric (for nearest neighbor selection) is sometimes critical
 - The local planner may be non-trivial

References


Steven M. LaValle: *Planning Algorithms*,
<http://planning.cs.uiuc.edu/>.

Choset et. al.: *Principles of Motion Planning*, MIT Press.


Latombe's "motion planning" lecture, <http://robotics.stanford.edu/~latombe/cs326/2007/schedule.htm>

RRT*


Sertac Karaman & Emilio Frazzoli: Incremental sampling-based algorithms for optimal motion planning, arXiv 1005.0416 (2010).




(a)




(b)




(c)




(d)




(e)



(f)



(g)



(h)

RRT*

Sertac Karaman & Emilio Frazzoli: Incremental sampling-based algorithms for optimal motion planning, arXiv 1005.0416 (2010).

Algorithm 4: $\text{Extend}_{\text{RRT}^*}(G, x)$

```
1  $V' \leftarrow V; E' \leftarrow E;$ 
2  $x_{\text{nearest}} \leftarrow \text{Nearest}(G, x);$ 
3  $x_{\text{new}} \leftarrow \text{Steer}(x_{\text{nearest}}, x);$ 
4 if  $\text{ObstacleFree}(x_{\text{nearest}}, x_{\text{new}})$  then
5    $V' \leftarrow V' \cup \{x_{\text{new}}\};$ 
6    $x_{\text{min}} \leftarrow x_{\text{nearest}};$ 
7    $X_{\text{near}} \leftarrow \text{Near}(G, x_{\text{new}}, |V|);$ 
8   for all  $x_{\text{near}} \in X_{\text{near}}$  do
9     if  $\text{ObstacleFree}(x_{\text{near}}, x_{\text{new}})$  then
10        $c' \leftarrow \text{Cost}(x_{\text{near}}) + c(\text{Line}(x_{\text{near}}, x_{\text{new}}));$ 
11       if  $c' < \text{Cost}(x_{\text{new}})$  then
12          $x_{\text{min}} \leftarrow x_{\text{near}};$ 
13    $E' \leftarrow E' \cup \{(x_{\text{min}}, x_{\text{new}})\};$ 
14   for all  $x_{\text{near}} \in X_{\text{near}} \setminus \{x_{\text{min}}\}$  do
15     if  $\text{ObstacleFree}(x_{\text{new}}, x_{\text{near}})$  and
16        $\text{Cost}(x_{\text{near}}) > \text{Cost}(x_{\text{new}}) + c(\text{Line}(x_{\text{new}}, x_{\text{near}}))$ 
17       then
18        $x_{\text{parent}} \leftarrow \text{Parent}(x_{\text{near}});$ 
19        $E' \leftarrow E' \setminus \{(x_{\text{parent}}, x_{\text{near}})\};$ 
20        $E' \leftarrow E' \cup \{(x_{\text{new}}, x_{\text{near}})\};$ 
21 return  $G' = (V', E')$ 
```

Non-holonomic systems

Non-holonomic systems

- We define a **nonholonomic system** as one with **differential constraints**:

$$\dim(u_t) < \dim(x_t)$$

\Rightarrow *Not all degrees of freedom are directly controllable*


- Dynamic systems are an example!
- General definition of a differential constraint:
For any given state x , let U_x be the tangent space that is generated by controls:

$$U_x = \{\dot{x} : \dot{x} = f(x, u), u \in U\}$$

$$(\text{non-holonomic} \iff \dim(U_x) < \dim(x))$$

The elements of U_x are elements of T_x subject to additional *differential constraints*.

Car example



$$\dot{x} = v \cos \theta$$

$$\dot{y} = v \sin \theta$$

$$\dot{\theta} = (v/L) \tan \varphi$$

$$|\varphi| < \Phi$$

State $q = \begin{pmatrix} x \\ y \\ \theta \end{pmatrix}$

System equation

Controls $u = \begin{pmatrix} v \\ \varphi \end{pmatrix}$

$$\begin{pmatrix} \dot{x} \\ \dot{y} \\ \dot{\theta} \end{pmatrix} = \begin{pmatrix} v \cos \theta \\ v \sin \theta \\ (v/L) \tan \varphi \end{pmatrix}$$

Car example

- The car is a *non-holonomic* system: not all DoFs are controlled, $\dim(u) < \dim(q)$
We have the *differential constraint* \dot{q} :


$$\dot{x} \sin \theta - \dot{y} \cos \theta = 0$$

“A car cannot move directly lateral.”

- Analogy to dynamic systems: Just like a car cannot instantly move sideways, a dynamic system cannot instantly change its position q : the current change in position is *constrained* by the current velocity \dot{q} .

Path finding for a non-holonomic system


Could a car follow this trajectory?



This is generated with a PRM in the state space $q = (x, y, \theta)$ *ignoring the differential constraint*.

Path finding with a non-holonomic system

This is a solution we would like to have:




The path respects **differential constraints**.

Each step in the path corresponds to setting certain controls.


Control-based sampling to grow a tree

- Control-based sampling: fulfils none of the nice exploration properties of RRTs, but fulfils the differential constraints:
 - 1) Select a $q \in T$ from tree of current configurations
 - 2) Pick control vector u at random
 - 3) Integrate equation of motion over short duration (picked at random or not)
 - 4) If the motion is collision-free, add the endpoint to the tree

Control-based sampling for the car




- 1) Select a $q \in T$
- 2) Pick v , ϕ , and τ
- 3) Integrate motion from q
- 4) Add result if collision-free




J. Barraquand and J.C. Latombe. Nonholonomic Multibody Robots:
Controllability and Motion Planning in the Presence of Obstacles. *Algorithmica*,
10:121-155, 1993.


car parking



car parking



parking with only left-steering



with a trailer

Better control-based exploration: RRTs revisited

- RRTs with differential constraints:

Input: q_{start} , number k of nodes, **time interval** τ

Output: tree $T = (V, E)$

1: initialize $V = \{q_{\text{start}}\}$, $E = \emptyset$

2: **for** $i = 0 : k$ **do**

3: $q_{\text{target}} \leftarrow$ random sample from Q

4: $q_{\text{near}} \leftarrow$ **nearest** neighbor of q_{target} in V

5: **use local planner to compute controls** u **that steer** q_{near} **towards** q_{target}

6: $q_{\text{new}} \leftarrow q_{\text{near}} + \int_{t=0}^{\tau} \dot{q}(q, u) dt$

7: **if** $q_{\text{new}} \in Q_{\text{free}}$ **then** $V \leftarrow V \cup \{q_{\text{new}}\}$, $E \leftarrow E \cup \{(q_{\text{near}}, q_{\text{new}})\}$

8: **end for**

- Crucial questions:
- How measure *near* in nonholonomic systems?
- How find controls u to steer towards target?

Configuration state metrics

Standard/Naive metrics:

- Comparing two 2D rotations/orientations $\theta_1, \theta_2 \in SO(2)$:
 - a) Euclidean metric between $e^{i\theta_1}$ and $e^{i\theta_2}$
 - b) $d(\theta_1, \theta_2) = \min\{|\theta_1 - \theta_2|, 2\pi - |\theta_1 - \theta_2|\}$
- Comparing two configurations $(x, y, \theta)_{1,2}$ in \mathbb{R}^2 :
Euclidean metric on $(x, y, e^{i\theta})$
- Comparing two 3D rotations/orientations $r_1, r_2 \in SO(3)$:
Represent both orientations as unit-length quaternions $r_1, r_2 \in \mathbb{R}^4$:
$$d(r_1, r_2) = \min\{|r_1 - r_2|, |r_1 + r_2|\}$$

where $|\cdot|$ is the Euclidean metric.
(Recall that r_1 and $-r_1$ represent exactly the same rotation.)

Configuration state metrics

Standard/Naive metrics:

- Comparing two 2D rotations/orientations $\theta_1, \theta_2 \in SO(2)$:
 - a) Euclidean metric between $e^{i\theta_1}$ and $e^{i\theta_2}$
 - b) $d(\theta_1, \theta_2) = \min\{|\theta_1 - \theta_2|, 2\pi - |\theta_1 - \theta_2|\}$
- Comparing two configurations $(x, y, \theta)_{1,2}$ in \mathbb{R}^2 :
Euclidean metric on $(x, y, e^{i\theta})$
- Comparing two 3D rotations/orientations $r_1, r_2 \in SO(3)$:
Represent both orientations as unit-length quaternions $r_1, r_2 \in \mathbb{R}^4$:
$$d(r_1, r_2) = \min\{|r_1 - r_2|, |r_1 + r_2|\}$$

where $|\cdot|$ is the Euclidean metric.
(Recall that r_1 and $-r_1$ represent exactly the same rotation.)

- **Ideal metric:**

Optimal cost-to-go between two states x_1 and x_2 :

- Use optimal trajectory cost as metric
- This is as hard to compute as the original problem, of course!!
→ Approximate, e.g., by neglecting obstacles.

Side story: Dubins curves

- Dubins car: constant velocity, and steer $\varphi \in [-\Phi, \Phi]$
- Neglecting obstacles, there are only **six** types of trajectories that connect any configuration $\in \mathbb{R}^2 \times \mathbb{S}^1$:


$$\{LRL, RLR, LSL, LSR, RSL, RSR\}$$

- annotating durations of each phase:

$$\{L_\alpha R_\beta L_\gamma, R_\alpha L_\beta R_\gamma, L_\alpha S_d L_\gamma, L_\alpha S_d R_\gamma, R_\alpha S_d L_\gamma, R_\alpha S_d R_\gamma\}$$


with $\alpha \in [0, 2\pi), \beta \in (\pi, 2\pi), d \geq 0$

Side story: Dubins curves



→ By testing all six types of trajectories for (q_1, q_2) we can define a Dubins metric for the RRT – and use the Dubins curves as controls!

Side story: Dubins curves



→ By testing all six types of trajectories for (q_1, q_2) we can define a Dubins metric for the RRT – and use the Dubins curves as controls!

- **Reeds-Shepp curves** are an extension for cars which can drive back. (includes 46 types of trajectories, good metric for use in RRTs for cars)