



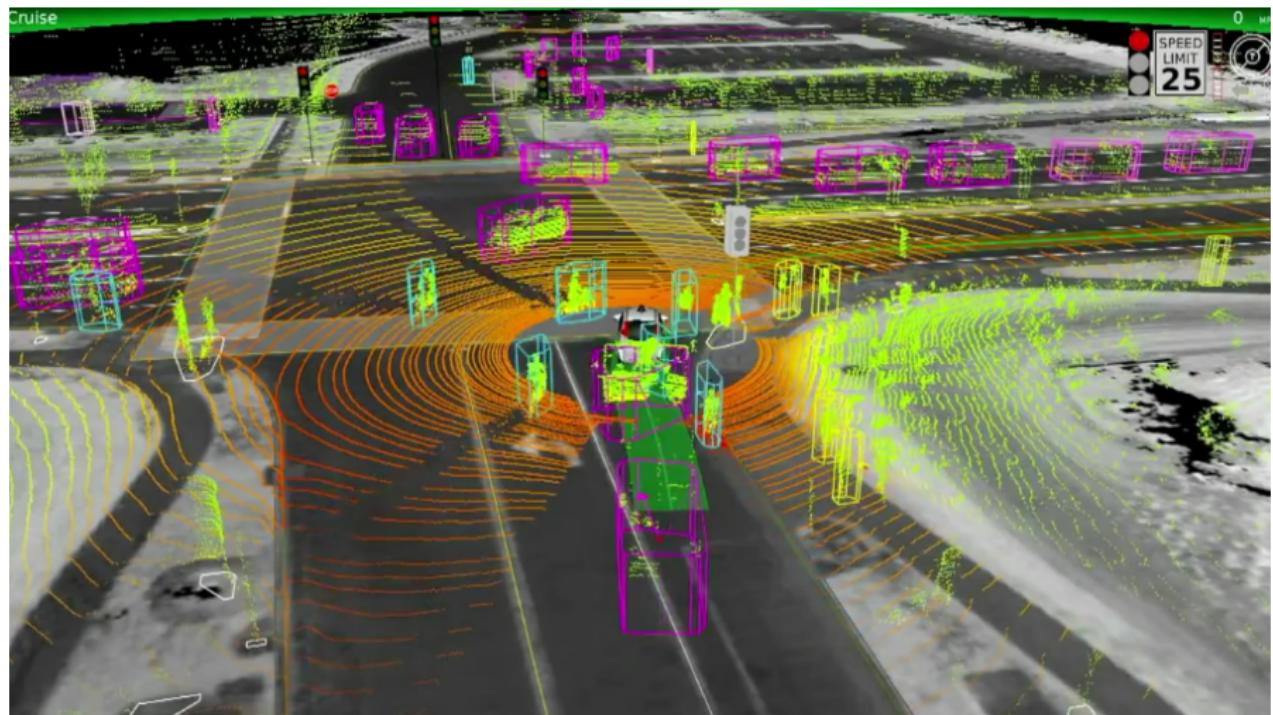
# Robotics

## Mobile Robotics

*State estimation, Bayes filter, odometry,  
particle filter, Kalman filter, SLAM, joint Bayes  
filter, EKF SLAM, particle SLAM*

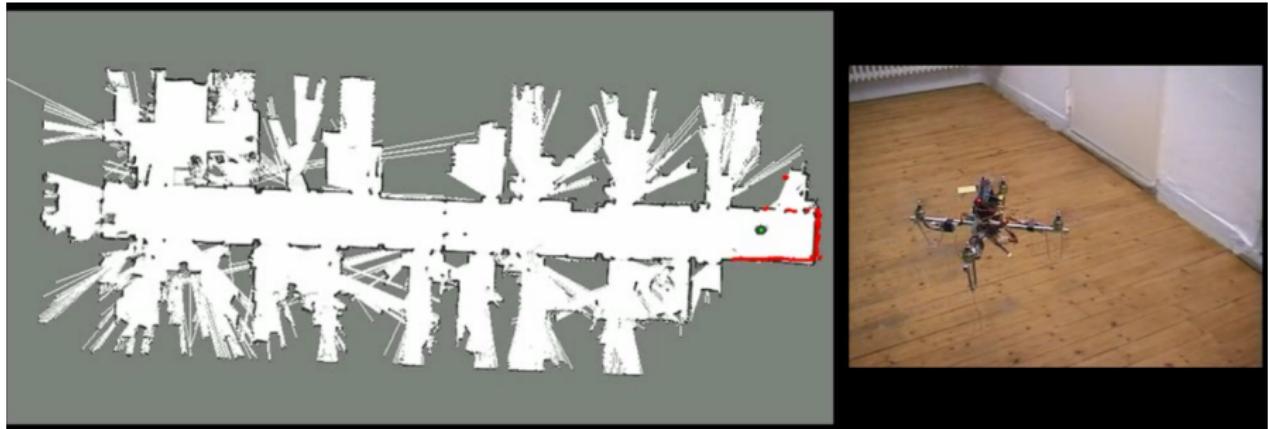
Marc Toussaint  
University of Stuttgart  
Winter 2016/17

Lecturer: Peter Englert



<https://www.youtube.com/watch?v=tiwVMrTLUWg>

Google Autonomous Car



<http://www.slawomir.de/publications/grzonka09icra/grzonka09icra.pdf>

## Quadcopter Indoor Localization

6x



<http://stair.stanford.edu/multimedia.php>

STAIR: STanford Artificial Intelligence Robot

# Outline

- PART I:
  - A core challenge in mobile robotics is **state estimation**
    - Bayesian filtering & smoothing
      - particle filter, Kalman filter

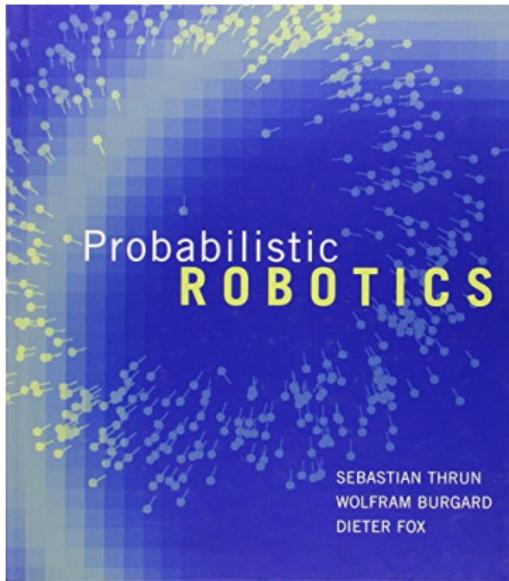
# Outline

- PART I:
  - A core challenge in mobile robotics is **state estimation**
    - Bayesian filtering & smoothing
      - particle filter, Kalman filter
- PART II:
  - Another challenge is to **build a map** while exploring
    - SLAM (simultaneous localization and mapping)

# Outline

- PART I:
  - A core challenge in mobile robotics is **state estimation**
    - Bayesian filtering & smoothing
      - particle filter, Kalman filter
- PART II:
  - Another challenge is to **build a map** while exploring
    - SLAM (simultaneous localization and mapping)

How to represent our knowledge? → **Probabilities**



*Sebastian Thrun, Wolfram Burghard,  
Dieter Fox: Probabilistic Robotics*  
[probabilistic-robotics.org](http://probabilistic-robotics.org)

# Types of Robot Mobility



- Each type of robot mobility corresponds to a system equation  $x_t = x_{t-1} + \tau f(x_{t-1}, u_{t-1})$
- or, if the dynamics are stochastic,

$$P(x_t | x_{t-1}, u_{t-1}) = \mathcal{N}(x_t | x_{t-1} + \tau f(x_{t-1}, u_{t-1}), \Sigma)$$

- Each type of robot mobility corresponds to a system equation  $x_t = x_{t-1} + \tau f(x_{t-1}, u_{t-1})$
- or, if the dynamics are stochastic,

$$P(x_t | x_{t-1}, u_{t-1}) = \mathcal{N}(x_t | x_{t-1} + \tau f(x_{t-1}, u_{t-1}), \Sigma)$$

- We considered control, path finding, and trajectory optimization

For this we always assumed to know the state  $x_t$  of the robot (e.g., its posture/position)!

- Each type of robot mobility corresponds to a system equation  $x_t = x_{t-1} + \tau f(x_{t-1}, u_{t-1})$   
or, if the dynamics are stochastic,

$$P(x_t | x_{t-1}, u_{t-1}) = \mathcal{N}(x_t | x_{t-1} + \tau f(x_{t-1}, u_{t-1}), \Sigma)$$

- We considered control, path finding, and trajectory optimization

For this we always assumed to know the state  $x_t$  of the robot (e.g., its posture/position)!

- Now we assume to observe sensor readings  $y_t$  (e.g., laser scan, image) and to have a sensor/observation model  $P(y_t | x_t)$ .

# **State Estimation**

# State Estimation Problem

- Our sensory data does not provide sufficient information to determine our location.
- Given the local sensor readings  $y_t$ , the current state  $x_t$  (location, position) is *uncertain*.
  - which hallway?
  - which door exactly?
  - which heading direction?



# State Estimation Problem

- What is the probability of being in front of room 154, given we see what is shown in the image?
- What is the probability given that we were just in front of room 156?
- What is the probability given that we were in front of room 156 and moved 15 meters?



## Recall Bayes' theorem

$$P(X|Y) = \frac{P(Y|X) P(X)}{P(Y)}$$

$$\text{posterior} = \frac{\text{likelihood} \cdot \text{prior}}{(\text{normalization})}$$

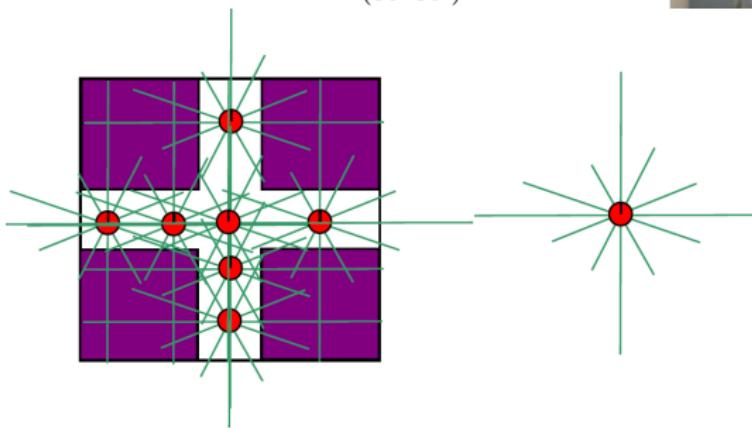
- How can we apply this to the State Estimation Problem?



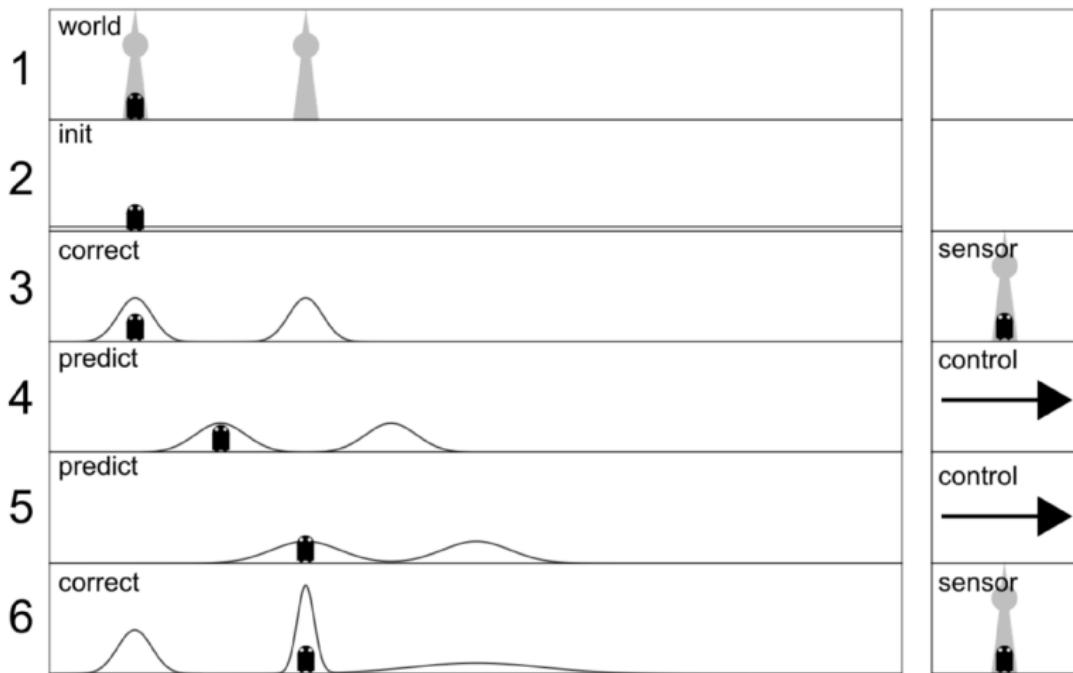
- How can we apply this to the State Estimation Problem?

Using Bayes Rule:

$$P(\text{location} \mid \text{sensor}) = \frac{P(\text{sensor} \mid \text{location})P(\text{location})}{P(\text{sensor})}$$



# Bayes Filter: 1d example



# Bayes Filter

$x_t$  = state (location) at time  $t$

$y_t$  = sensor readings at time  $t$

$u_{t-1}$  = control command (action, steering, velocity) at time  $t-1$

transition model  $P(x_t | u_{t-1}, x_{t-1})$

observation model  $P(y_t | x_t)$

- Given the history  $y_{0:t}$  and  $u_{0:t-1}$ , we want to compute the probability distribution over the state at time  $t$

$$p_t(x_t) := P(x_t | y_{0:t}, u_{0:t-1})$$

# Bayes Filter

$$p_t(x_t) := P(x_t \mid y_{0:t}, u_{0:t-1})$$

# Bayes Filter

$$\begin{aligned} p_t(x_t) &:= P(x_t \mid y_{0:t}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t, y_{0:t-1}, u_{0:t-1}) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \end{aligned}$$

using Bayes rule  $P(X|Y, Z) = c P(Y|X, Z) P(X|Z)$  with some normalization constant  $c_t$

# Bayes Filter

$$\begin{aligned} p_t(x_t) &:= P(x_t \mid y_{0:t}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t, y_{0:t-1}, u_{0:t-1}) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \end{aligned}$$

uses conditional independence of the observation on past observations  
and controls

# Bayes Filter

$$\begin{aligned} p_t(x_t) &:= P(x_t \mid y_{0:t}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t, y_{0:t-1}, u_{0:t-1}) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t, x_{t-1} \mid y_{0:t-1}, u_{0:t-1}) dx_{t-1} \end{aligned}$$

by definition of the marginal

# Bayes Filter

$$\begin{aligned} p_t(x_t) &:= P(x_t \mid y_{0:t}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t, y_{0:t-1}, u_{0:t-1}) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t, x_{t-1} \mid y_{0:t-1}, u_{0:t-1}) dx_{t-1} \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t \mid x_{t-1}, y_{0:t-1}, u_{0:t-1}) P(x_{t-1} \mid y_{0:t-1}, u_{0:t-1}) dx_{t-1} \end{aligned}$$

by definition of a conditional

# Bayes Filter

$$\begin{aligned} p_t(x_t) &:= P(x_t \mid y_{0:t}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t, y_{0:t-1}, u_{0:t-1}) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t, x_{t-1} \mid y_{0:t-1}, u_{0:t-1}) dx_{t-1} \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t \mid x_{t-1}, y_{0:t-1}, u_{0:t-1}) P(x_{t-1} \mid y_{0:t-1}, u_{0:t-1}) dx_{t-1} \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t \mid x_{t-1}, u_{t-1}) P(x_{t-1} \mid y_{0:t-1}, u_{0:t-2}) dx_{t-1} \end{aligned}$$

given  $x_{t-1}$ ,  $x_t$  depends only on the controls  $u_{t-1}$  (Markov Property)

# Bayes Filter

$$\begin{aligned} p_t(x_t) &:= P(x_t \mid y_{0:t}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t, y_{0:t-1}, u_{0:t-1}) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) P(x_t \mid y_{0:t-1}, u_{0:t-1}) \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t, x_{t-1} \mid y_{0:t-1}, u_{0:t-1}) dx_{t-1} \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t \mid x_{t-1}, y_{0:t-1}, u_{0:t-1}) P(x_{t-1} \mid y_{0:t-1}, u_{0:t-1}) dx_{t-1} \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t \mid x_{t-1}, u_{t-1}) P(x_{t-1} \mid y_{0:t-1}, u_{0:t-2}) dx_{t-1} \\ &= c_t P(y_t \mid x_t) \int_{x_{t-1}} P(x_t \mid u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1} \end{aligned}$$

- A Bayes filter updates the posterior belief  $p_t(x_t)$  in each time step using the:

observation model  $P(y_t \mid x_t)$

transition model  $P(x_t \mid u_{t-1}, x_{t-1})$

old estimate  $p_{t-1}(x_{t-1})$

# Bayes Filter

$$p_t(x_t) \propto \underbrace{P(y_t | x_t)}_{\text{new information}} \underbrace{\int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) \underbrace{p_{t-1}(x_{t-1})}_{\text{old estimate}} dx_{t-1}}_{\text{predictive estimate } \hat{p}_t(x_t)}$$

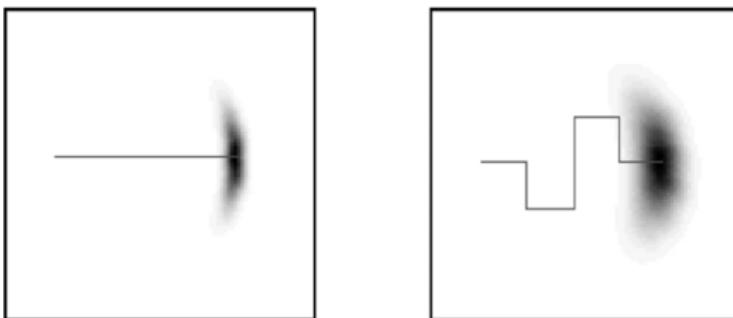
1. We have a belief  $p_{t-1}(x_{t-1})$  of our previous position
2. We use the motion model to predict the current position

$$\hat{p}_t(x_t) \propto \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1}$$

3. We integrate this with the current observation to get a better belief

$$p_t(x_t) \propto P(y_t | x_t) \hat{p}_t(x_t)$$

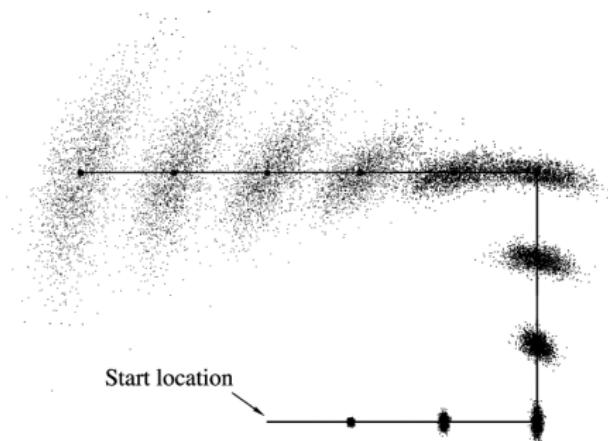
- Typical transition model  $P(x_t | u_{t-1}, x_{t-1})$  in robotics:



(from *Robust Monte Carlo localization for mobile robots* Sebastian Thrun, Dieter Fox, Wolfram Burgard, Frank Dellaert)

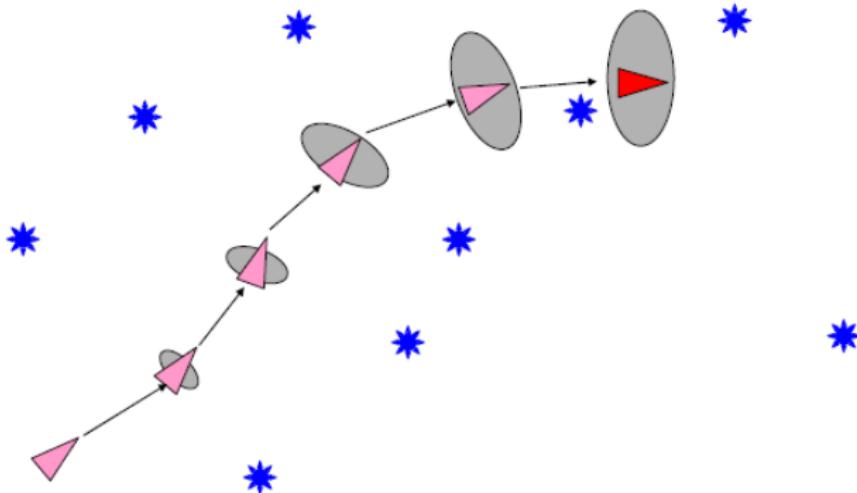
# Odometry (“Dead Reckoning”): Filtering without observations

- The predictive distributions  $\hat{p}_t(x_t)$  *without integrating observations* (removing the  $P(y_t|x_t)$  part from the Bayesian filter)

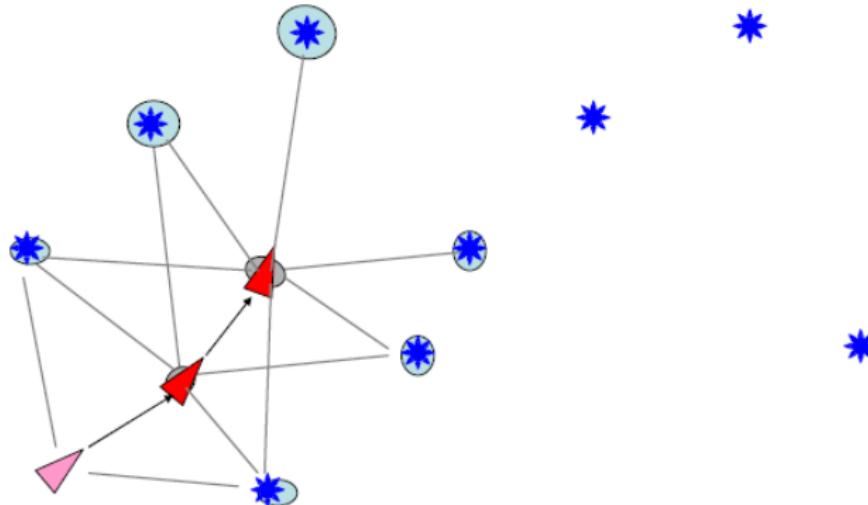


(from *Robust Monte Carlo localization for mobile robots* Sebastian Thrun, Dieter Fox, Wolfram Burgard, Frank Dellaert)

Again, predictive distributions  $\hat{p}_t(x_t)$  *without integrating landmark observations*



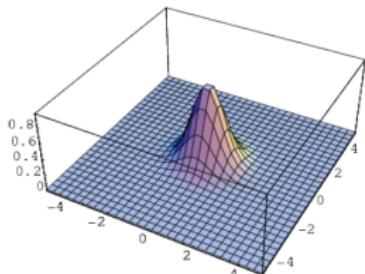
The Bayes-filtered distributions  $p_t(x_t)$  integrating landmark observations



# Bayesian Filters

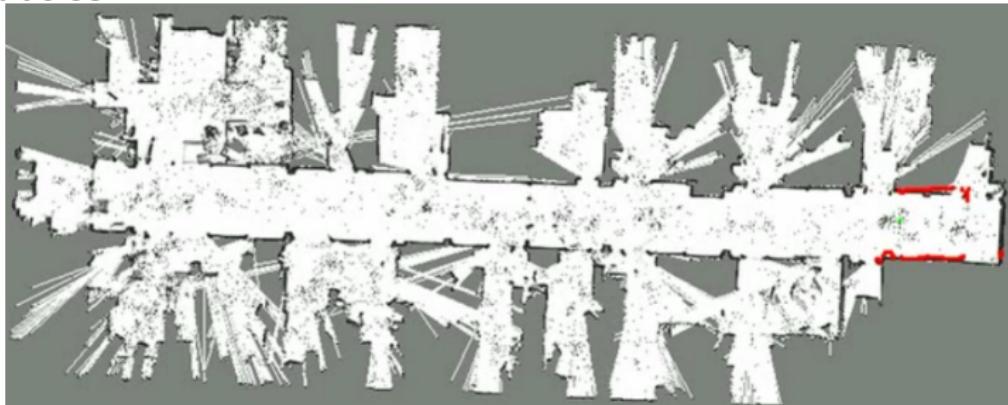
- How to represent the belief  $p_t(x_t)$ :

- Gaussian



$$\frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

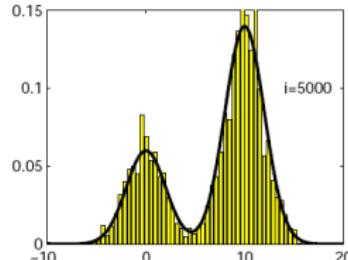
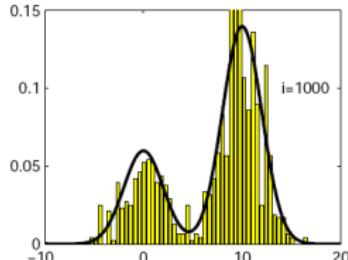
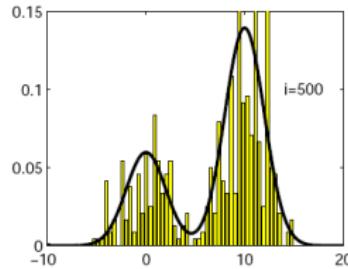
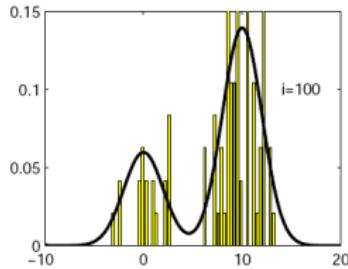
- Particles



# Recall: Particle Representation of a Distribution

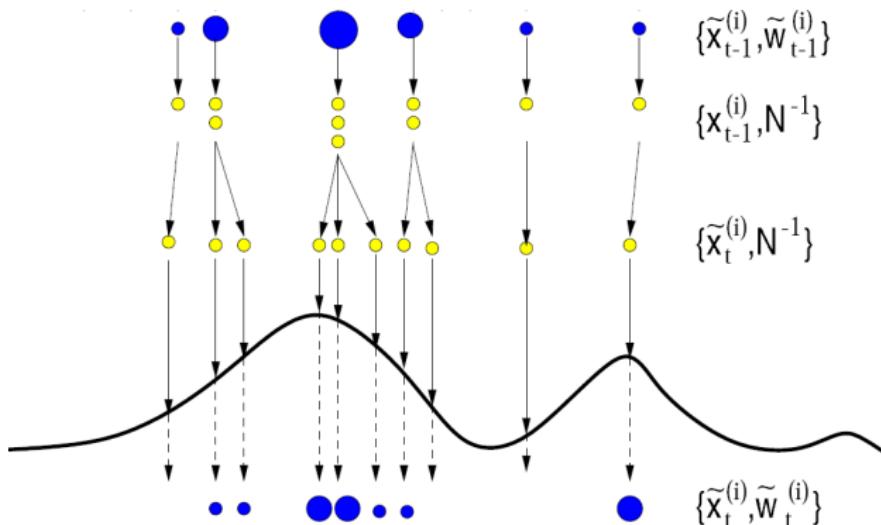
- Weighed set of  $N$  particles  $\{(x^i, w^i)\}_{i=1}^N$

$$p(x) \approx q(x) := \sum_{i=1}^N w^i \delta(x, x^i)$$



# Particle Filter := Bayesian Filtering with Particles

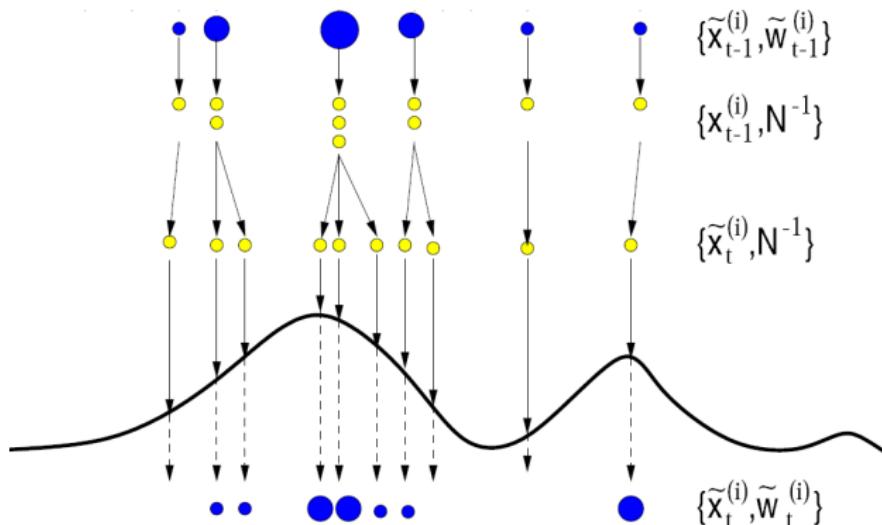
$$(\text{Bayes Filter: } p_t(x_t) \propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1})$$



1. Start with  $N$  particles  $\{(x_{t-1}^i, w_{t-1}^i)\}_{i=1}^N$

# Particle Filter := Bayesian Filtering with Particles

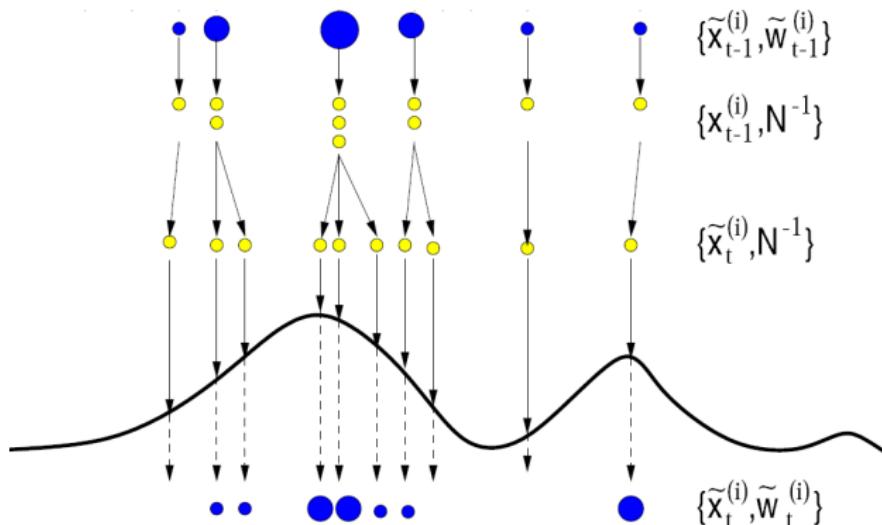
$$(\text{Bayes Filter: } p_t(x_t) \propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1})$$



1. Start with  $N$  particles  $\{(x_{t-1}^i, w_{t-1}^i)\}_{i=1}^N$
2. Resample particles to get  $N$  weight-1-particles:  $\{\hat{x}_{t-1}^i\}_{i=1}^N$

# Particle Filter := Bayesian Filtering with Particles

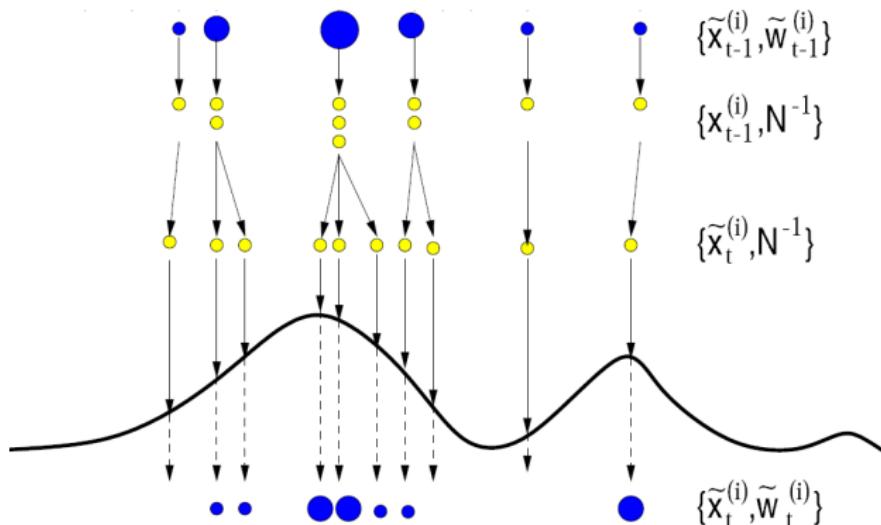
$$(\text{Bayes Filter: } p_t(x_t) \propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1})$$



1. Start with  $N$  particles  $\{(x_{t-1}^i, w_{t-1}^i)\}_{i=1}^N$
2. Resample particles to get  $N$  weight-1-particles:  $\{\hat{x}_{t-1}^i\}_{i=1}^N$
3. Use motion model to get new “predictive” particles  $\{x_t^i\}_{i=1}^N$   
for each  $x_t^i \sim P(x_t | u_{t-1}, \hat{x}_{t-1}^i)$

# Particle Filter := Bayesian Filtering with Particles

$$(\text{Bayes Filter: } p_t(x_t) \propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1})$$



1. Start with  $N$  particles  $\{(x_{t-1}^i, w_{t-1}^i)\}_{i=1}^N$
2. Resample particles to get  $N$  weight-1-particles:  $\{\hat{x}_{t-1}^i\}_{i=1}^N$
3. Use motion model to get new “predictive” particles  $\{x_t^i\}_{i=1}^N$   
for each  $x_t^i \sim P(x_t | u_{t-1}, \hat{x}_{t-1}^i)$
4. Use observation model to assign new weights  $w_t^i \propto P(y_t | x_t^i)$

# Practical Considerations

- Given a complex multi-modal distribution: *where are we?*
  - particle with the highest weight?
  - which *cluster*?
  - kernel density estimation: cluster with highest density?

# Practical Considerations

- Given a complex multi-modal distribution: *where are we?*
  - particle with the highest weight?
  - which *cluster*?
  - kernel density estimation: cluster with highest density?
- Particle deprivation:  $\{x_t^i\}_{i=1}^N$  have very low observation likelihood  $P(y_t | x_t^i)$ : “particle die over time”
  - random samples
  - *sensor reset*: samples from  $P(y_t | x_t)$
  - ...

# Practical Considerations

- Given a complex multi-modal distribution: *where are we?*
  - particle with the highest weight?
  - which *cluster*?
  - kernel density estimation: cluster with highest density?
- Particle deprivation:  $\{x_t^i\}_{i=1}^N$  have very low observation likelihood  $P(y_t | x_t^i)$ : “particle die over time”
  - random samples
  - *sensor reset*: samples from  $P(y_t | x_t)$
  - ...
- How many particles? → adaptive techniques

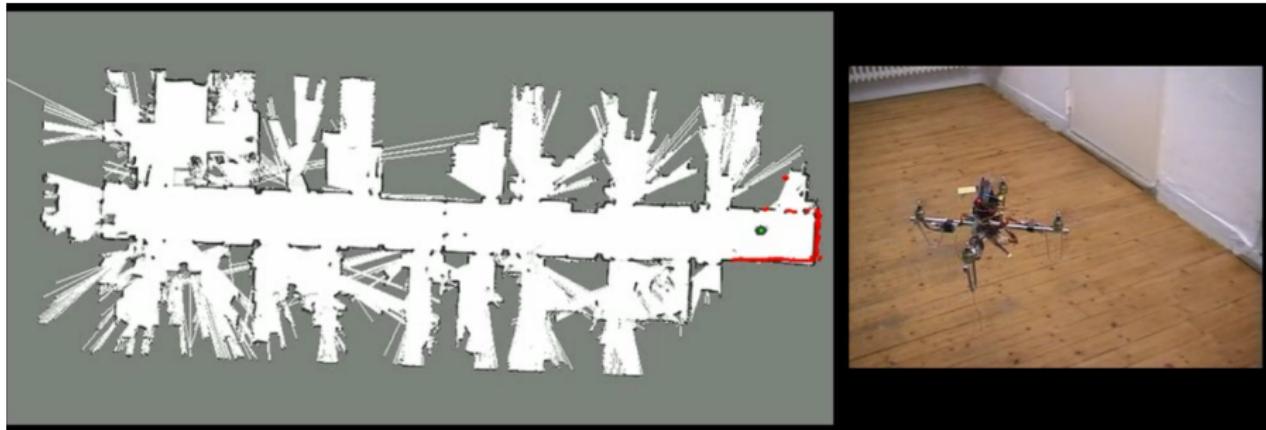
# Practical Considerations

- Given a complex multi-modal distribution: *where are we?*
  - particle with the highest weight?
  - which *cluster*?
  - kernel density estimation: cluster with highest density?
- Particle deprivation:  $\{x_t^i\}_{i=1}^N$  have very low observation likelihood  $P(y_t | x_t^i)$ : “particle die over time”
  - random samples
  - *sensor reset*: samples from  $P(y_t | x_t)$
  - ...
- How many particles? → adaptive techniques
- Resampling is important: many variations exist!  
*Liu & Chen (1998): Sequential Monte Carlo Methods for Dynamic Systems.*  
*Douc, Cappé & Moulines: Comparison of Resampling Schemes for Particle Filtering.*

# Particle Filter: Summary

- conceptually simple
  - approximates the true distribution  $p(x)$
  - non-parametric
  - can deal with non-linear transformations
- 
- “Particle Filter” in different communities
    - *Monte Carlo Localization* in the mobile robotics community
    - *Condensation Algorithm* in the vision community

## Example: Quadcopter Localization



<http://www.slawomir.de/publications/grzonka09icra/grzonka09icra.pdf>  
Quadcopter Indoor Localization

# Kalman filter := Bayesian Filtering with Gaussians

Bayes Filter:  $p_t(x_t) \propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1}$

- Can be computed analytically for linear-Gaussian observations and transitions:

$$P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$$

$$P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | A(u_{t-1}) x_{t-1} + a(u_{t-1}), Q)$$

Definition:

$$\mathcal{N}(x | a, A) = \frac{1}{|2\pi A|^{1/2}} \exp\left\{-\frac{1}{2}(x - a)^\top A^{-1} (x - a)\right\}$$

Product:

$$\mathcal{N}(x | a, A) \mathcal{N}(x | b, B) = \mathcal{N}(x | B(A+B)^{-1}a + A(A+B)^{-1}b, A(A+B)^{-1}B) \mathcal{N}(a | b, A + B)$$

“Propagation”:

$$\int_y \mathcal{N}(x | a + Fy, A) \mathcal{N}(y | b, B) dy = \mathcal{N}(x | a + Fb, A + FBF^\top)$$

Transformation:

$$\mathcal{N}(Fx + f | a, A) = \frac{1}{|F|} \mathcal{N}(x | F^{-1}(a - f), F^{-1}AF^{-\top})$$

(more identities: see “Gaussian identities”

<http://ipvs.informatik.uni-stuttgart.de/mlr/marc/notes/gaussians.pdf>)

# Kalman filter derivation

state:  $p_t(x_t) = \mathcal{N}(x_t | s_t, S_t)$

transition model:  $P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | Ax_{t-1} + a, Q)$

observation model:  $P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$

$$p_t(x_t) \propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1}$$

# Kalman filter derivation

state:  $p_t(x_t) = \mathcal{N}(x_t | s_t, S_t)$

transition model:  $P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | Ax_{t-1} + a, Q)$

observation model:  $P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$

$$\begin{aligned} p_t(x_t) &\propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \int_{x_{t-1}} \mathcal{N}(x_t | Ax_{t-1} + a, Q) \mathcal{N}(x_{t-1} | s_{t-1}, S_{t-1}) dx_{t-1} \end{aligned}$$

# Kalman filter derivation

state:  $p_t(x_t) = \mathcal{N}(x_t | s_t, S_t)$

transition model:  $P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | Ax_{t-1} + a, Q)$

observation model:  $P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$

$$\begin{aligned} p_t(x_t) &\propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \int_{x_{t-1}} \mathcal{N}(x_t | Ax_{t-1} + a, Q) \mathcal{N}(x_{t-1} | s_{t-1}, S_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \mathcal{N}(x_t | \underbrace{As_{t-1} + a}_{=: \hat{s}_t}, \underbrace{Q + AS_{t-1}A^\top}_{=: \hat{S}_t}) \end{aligned}$$

“Propagation”:  $\int_y \mathcal{N}(x | a + Fy, A) \mathcal{N}(y | b, B) dy = \mathcal{N}(x | a + Fb, A + FBF^\top)$

# Kalman filter derivation

state:  $p_t(x_t) = \mathcal{N}(x_t | s_t, S_t)$

transition model:  $P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | Ax_{t-1} + a, Q)$

observation model:  $P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$

$$\begin{aligned} p_t(x_t) &\propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \int_{x_{t-1}} \mathcal{N}(x_t | Ax_{t-1} + a, Q) \mathcal{N}(x_{t-1} | s_{t-1}, S_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \mathcal{N}(x_t | \underbrace{As_{t-1} + a}_{=: \hat{s}_t}, \underbrace{Q + AS_{t-1}A^\top}_{=: \hat{S}_t}) \\ &= \mathcal{N}(Cx_t + c | y_t, W) \mathcal{N}(x_t | \hat{s}_t, \hat{S}_t) \end{aligned}$$

Use symmetry of Gaussian:  $\mathcal{N}(x | a, A) = \mathcal{N}(a | x, A)$

# Kalman filter derivation

state:  $p_t(x_t) = \mathcal{N}(x_t | s_t, S_t)$

transition model:  $P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | Ax_{t-1} + a, Q)$

observation model:  $P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$

$$\begin{aligned} p_t(x_t) &\propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \int_{x_{t-1}} \mathcal{N}(x_t | Ax_{t-1} + a, Q) \mathcal{N}(x_{t-1} | s_{t-1}, S_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \mathcal{N}(x_t | \underbrace{As_{t-1} + a}_{=: \hat{s}_t}, \underbrace{Q + AS_{t-1}A^\top}_{=: \hat{S}_t}) \\ &= \mathcal{N}(Cx_t + c | y_t, W) \mathcal{N}(x_t | \hat{s}_t, \hat{S}_t) \\ &= \mathcal{N}(x_t | s_t, S_t) \cdot \langle \text{terms indep. of } x_t \rangle \end{aligned}$$

# Kalman filter derivation

state:  $p_t(x_t) = \mathcal{N}(x_t | s_t, S_t)$

transition model:  $P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | Ax_{t-1} + a, Q)$

observation model:  $P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$

$$\begin{aligned} p_t(x_t) &\propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \int_{x_{t-1}} \mathcal{N}(x_t | Ax_{t-1} + a, Q) \mathcal{N}(x_{t-1} | s_{t-1}, S_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \mathcal{N}(x_t | \underbrace{As_{t-1} + a}_{=: \hat{s}_t}, \underbrace{Q + AS_{t-1}A^\top}_{=: \hat{S}_t}) \\ &= \mathcal{N}(Cx_t + c | y_t, W) \mathcal{N}(x_t | \hat{s}_t, \hat{S}_t) \\ &= \mathcal{N}(x_t | s_t, S_t) \cdot \langle \text{terms indep. of } x_t \rangle \\ S_t &= (C^\top W^{-1} C + \hat{S}_t^{-1})^{-1} = \hat{S}_t - \underbrace{\hat{S}_t C^\top (W + C \hat{S}_t C^\top)^{-1} C \hat{S}_t}_{\text{"Kalman gain" } K} \end{aligned}$$

new covariance  $S_t$ ; uses the general Woodbury identity

# Kalman filter derivation

state:  $p_t(x_t) = \mathcal{N}(x_t | s_t, S_t)$

transition model:  $P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | Ax_{t-1} + a, Q)$

observation model:  $P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$

$$\begin{aligned} p_t(x_t) &\propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \int_{x_{t-1}} \mathcal{N}(x_t | Ax_{t-1} + a, Q) \mathcal{N}(x_{t-1} | s_{t-1}, S_{t-1}) dx_{t-1} \\ &= \mathcal{N}(y_t | Cx_t + c, W) \mathcal{N}(x_t | \underbrace{As_{t-1} + a}_{=: \hat{s}_t}, \underbrace{Q + AS_{t-1}A^\top}_{=: \hat{S}_t}) \\ &= \mathcal{N}(Cx_t + c | y_t, W) \mathcal{N}(x_t | \hat{s}_t, \hat{S}_t) \\ &= \mathcal{N}(x_t | s_t, S_t) \cdot \langle \text{terms indep. of } x_t \rangle \\ S_t &= (C^\top W^{-1} C + \hat{S}_t^{-1})^{-1} = \hat{S}_t - \underbrace{\hat{S}_t C^\top (W + C \hat{S}_t C^\top)^{-1} C \hat{S}_t}_{\text{"Kalman gain" } K} \\ s_t &= S_t [C^\top W^{-1} (y_t - c) + \hat{S}_t^{-1} \hat{s}_t] = \hat{s}_t + K(y_t - C \hat{s}_t - c) \end{aligned}$$

new mean  $s_t$ ; uses  $S_t C^\top W^{-1} = K$  and  $S_t \hat{S}_t^{-1} = \mathbf{I} - KC$

# Extended Kalman filter (EKF) and Unscented Transform

Bayes Filter:  $p_t(x_t) \propto P(y_t | x_t) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}) dx_{t-1}$

- Can be computed analytically for linear-Gaussian observations and transitions:

$$P(y_t | x_t) = \mathcal{N}(y_t | Cx_t + c, W)$$

$$P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | A(u_{t-1})x_{t-1} + a(u_{t-1}), Q)$$

- If  $P(y_t | x_t)$  or  $P(x_t | u_{t-1}, x_{t-1})$  are not linear:

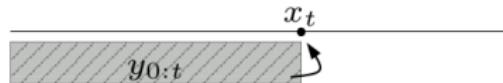
$$P(y_t | x_t) = \mathcal{N}(y_t | g(x_t), W)$$

$$P(x_t | u_{t-1}, x_{t-1}) = \mathcal{N}(x_t | f(x_{t-1}, u_{t-1}), Q)$$

- approximate  $f$  and  $g$  as locally linear (*Extended Kalman Filter*)
  - or sample locally from them and reapproximate as Gaussian (*Unscented Transform*)

# Bayes smoothing

Filtering:  $P(x_t | y_{0:t})$



Smoothing:  $P(x_t | y_{0:T})$



Prediction:  $P(x_t | y_{0:s})$



# Bayes smoothing

- Let  $\mathcal{P} = y_{0:t}$  past observations,  $\mathcal{F} = y_{t+1:T}$  future observations

$$\begin{aligned} P(x_t | \mathcal{P}, \mathcal{F}, u_{0:T}) &\propto P(\mathcal{F} | x_t, \mathcal{P}, u_{0:T}) P(x_t | \mathcal{P}, u_{0:T}) \\ &= \underbrace{P(\mathcal{F} | x_t, u_{t:T})}_{=: \beta_t(x_t)} \underbrace{P(x_t | \mathcal{P}, u_{0:t-1})}_{=: p(x_t)} \end{aligned}$$

*Bayesian smoothing fuses a forward filter  $p_t(x_t)$  with a backward “filter”  $\beta_t(x_t)$*

# Bayes smoothing

- Let  $\mathcal{P} = y_{0:t}$  past observations,  $\mathcal{F} = y_{t+1:T}$  future observations

$$\begin{aligned} P(x_t | \mathcal{P}, \mathcal{F}, u_{0:T}) &\propto P(\mathcal{F} | x_t, \mathcal{P}, u_{0:T}) P(x_t | \mathcal{P}, u_{0:T}) \\ &= \underbrace{P(\mathcal{F} | x_t, u_{t:T})}_{=: \beta_t(x_t)} \underbrace{P(x_t | \mathcal{P}, u_{0:t-1})}_{=: p(x_t)} \end{aligned}$$

*Bayesian smoothing fuses a forward filter  $p_t(x_t)$  with a backward “filter”  $\beta_t(x_t)$*

- Backward recursion (derivation analogous to the Bayesian filter)

$$\begin{aligned} \beta_t(x_t) &:= P(y_{t+1:T} | x_t, u_{t:T}) \\ &= \int_{x_{t+1}} \beta_{t+1}(x_{t+1}) P(y_{t+1} | x_{t+1}) P(x_{t+1} | x_t, u_t) dx_{t+1} \end{aligned}$$

# **Simultaneous Localization and Mapping (SLAM)**

## Localization and Mapping

- The Bayesian filter requires an observation model  $P(y_t | x_t)$
- A **map** is something that provides the observation model:  
*A map tells us for each  $x_t$  what the sensor readings  $y_t$  might look like*

# Types of maps

## Grid map



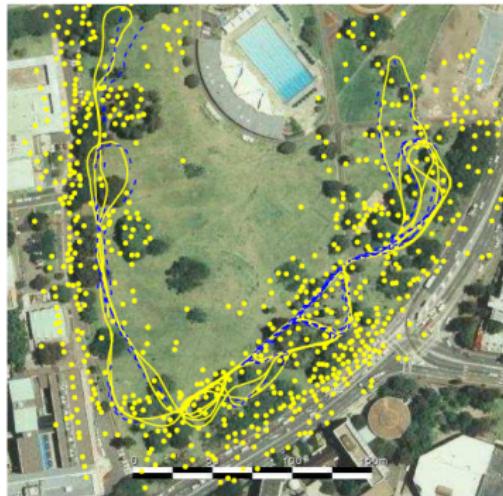
K. Murphy (1999): *Bayesian map learning in dynamic environments*.

Grisetti, Tipaldi, Stachniss, Burgard, Nardi: *Fast and Accurate SLAM with Rao-Blackwellized Particle Filters*

## Laser scan map



## Landmark map



## Victoria Park data set

M. Montemerlo, S. Thrun, D. Koller, & B. Wegbreit (2003): *FastSLAM 2.0: An improved particle filtering algorithm for simultaneous localization and mapping that provably converges*. IJCAI, 1151–1156.

# Simultaneous Localization and Mapping Problem

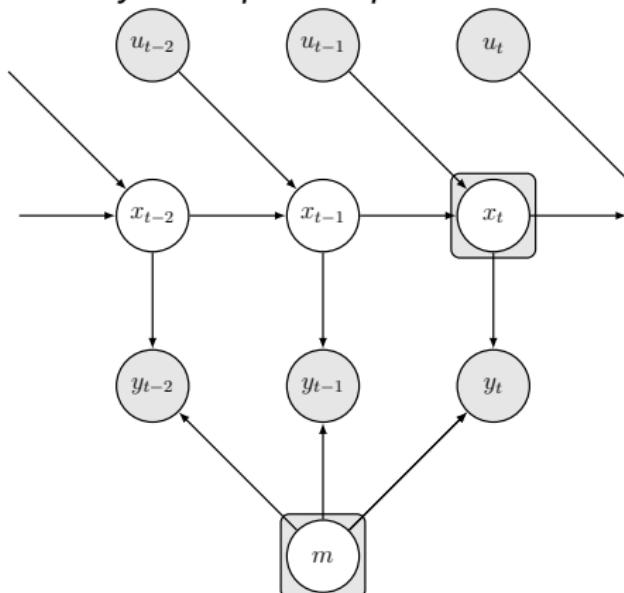
- Notation:

$x_t$  = state (location) at time  $t$

$y_t$  = sensor readings at time  $t$

$u_{t-1}$  = control command (action, steering, velocity) at time  $t-1$

$m$  = the map; formally: a map is the parameters that define  $P(y_t | x_t, m)$



# Simultaneous Localization and Mapping Problem

- Notation:

$x_t$  = state (location) at time  $t$

$y_t$  = sensor readings at time  $t$

$u_{t-1}$  = control command (action, steering, velocity) at time  $t-1$

$m$  = the map; *formally: a map is the parameters that define  $P(y_t | x_t, m)$*

- Given the history  $y_{0:t}$  and  $u_{0:t-1}$ , we want to compute the belief over the pose and the map  $m$  at time  $t$

$$p_t(x_t, m) := P(x_t, m | y_{0:t}, u_{0:t-1})$$

- We assume to know the:

– transition model  $P(x_t | u_{t-1}, x_{t-1})$

– observation model  $P(y_t | x_t, m)$  (*defined by the map*)

# SLAM: classical “chicken or egg problem”

- If we knew the map we could use a Bayes filter to compute the belief over the state

$$P(x_t | m, y_{0:t}, u_{0:t-1})$$

- If we knew the state trajectory  $x_{0:t}$  we could efficiently compute the belief over the map

$$P(m | x_{0:t}, y_{0:t}, u_{0:t-1})$$

# SLAM: classical “chicken or egg problem”

- If we knew the map we could use a Bayes filter to compute the belief over the state

$$P(x_t | m, y_{0:t}, u_{0:t-1})$$

- If we knew the state trajectory  $x_{0:t}$  we could efficiently compute the belief over the map

$$P(m | x_{0:t}, y_{0:t}, u_{0:t-1})$$

- SLAM requires to tie state estimation and map building together:

- 1) Joint inference on  $x_t$  and  $m$

→ Kalman-SLAM

- 2) Tie a state hypothesis (=particle) to a map hypothesis

→ particle SLAM

- 3) Frame everything as a graph optimization problem

→ graph SLAM

## Joint Bayesian Filter over $x$ and $m$

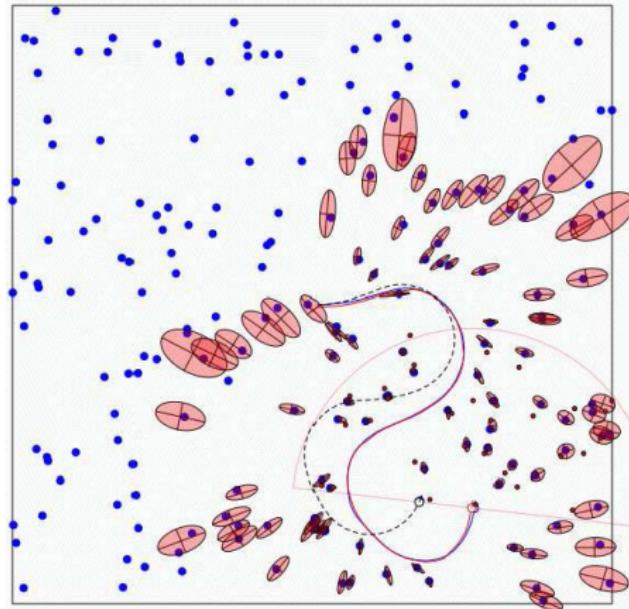
- A (formally) straight-forward approach is the joint Bayesian filter

$$p_t(x_t, m) \propto P(y_t | x_t, m) \int_{x_{t-1}} P(x_t | u_{t-1}, x_{t-1}) p_{t-1}(x_{t-1}, m) dx_{t-1}$$

But: How to represent a belief over high-dimensional  $x_t, m$ ?

# Map uncertainty

- In this case the map  $m = (\theta_1, \dots, \theta_N)$  is a set of  $N$  landmarks,  $\theta_j \in \mathbb{R}^2$

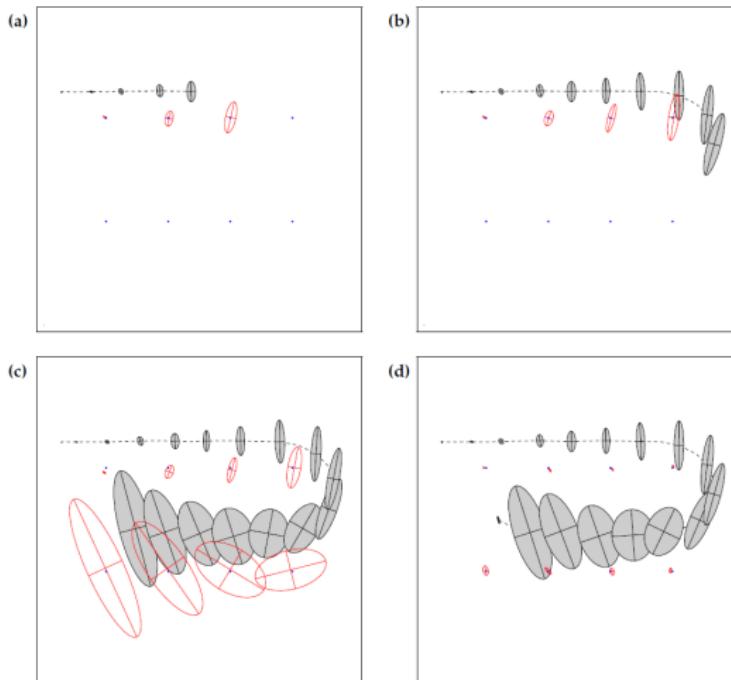


- Use Gaussians to represent the uncertainty of landmark positions

## (Extended) Kalman Filter SLAM

- Analogous to Localization with Gaussian for the pose belief  $p_t(x_t)$ 
  - But now: joint belief  $p_t(x_t, \theta_{1:N})$  is  $3 + 2N$ -dimensional Gaussian
  - Assumes the map  $m = (\theta_1, \dots, \theta_N)$  is a set of  $N$  landmarks,  $\theta_j \in \mathbb{R}^2$
  - Exact update equations (under the Gaussian assumption)
  - Conceptually very simple
- Drawbacks:
  - Scaling (full covariance matrix is  $O(N^2)$ )
  - Sometimes non-robust (uni-modal, “data association problem”)
  - Lacks advantages of Particle Filter
    - (multiple hypothesis, more robust to non-linearities)

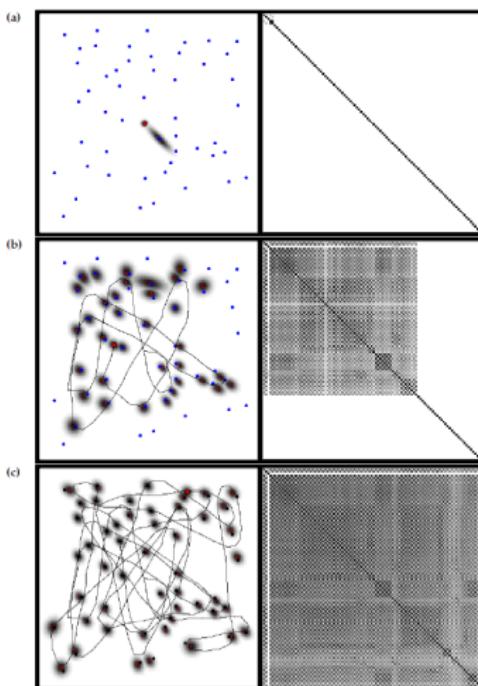
# EKF Slam Example 1



(from *Probabilistic Robotics* Sebastian Thrun, Wolfram Burgard, Dieter Fox)

- The landmark uncertainty decreases when the robot sees the first landmark again.

## EKF Slam Example 2



(from *Probabilistic Robotics* Sebastian Thrun, Wolfram Burgard, Dieter Fox)

- The right images show the correlation matrix. After some time all landmarks become fully correlated.

# SLAM with particles

**Core idea:** Each particle carries its own map belief

# SLAM with particles

**Core idea: Each particle carries its own map belief**

- As for the localization problem use particles to represent the *pose belief*  $p_t(x_t)$

(Note: Each particle actually “has a history  $x_{0:t}^i$ ” – a whole trajectory!)

- Factorization of the SLAM posterior:

$$P(x_{0:t}, m | y_{0:t}, u_{0:t-1}) = P(x_{0:t} | y_{0:t}, u_{0:t-1}) \underbrace{P(m | x_{0:t}, y_{0:t})}_{p_t(m)}$$

# SLAM with particles

**Core idea:** Each particle carries its own map belief

- As for the localization problem use particles to represent the *pose belief*  $p_t(x_t)$   
(Note: Each particle actually “has a history  $x_{0:t}^i$ ” – a whole trajectory!)
- Factorization of the SLAM posterior:  
$$P(x_{0:t}, m | y_{0:t}, u_{0:t-1}) = P(x_{0:t} | y_{0:t}, u_{0:t-1}) \underbrace{P(m | x_{0:t}, y_{0:t})}_{p_t(m)}$$
- Each particle  $i$  stores its path estimate  $x_{0:t}^i$  and an estimate of the map belief  $p_t^i(m)$  conditioned on the particle history  $x_{0:t}^i$ . The conditional beliefs  $p_t^i(m)$  may be factorized over grid points or landmarks of the map.

K. Murphy (1999): *Bayesian map learning in dynamic environments.*

[http://www.cs.ubc.ca/~murphyk/Papers/map\\_nips99.pdf](http://www.cs.ubc.ca/~murphyk/Papers/map_nips99.pdf)

## Map estimation for a *given* particle history

- Given  $x_{0:t}$  (e.g. a trajectory of a particle), what is the posterior over the map  $m$ ?

→ simplified Bayes Filter:

$$p_t(m) := P(m \mid x_{0:t}, y_{0:t}) \propto P(y_t \mid m, x_t) p_{t-1}(m)$$

(no transition model: assumption that map is constant)

# Map estimation for a *given* particle history

- Given  $x_{0:t}$  (e.g. a trajectory of a particle), what is the posterior over the map  $m$ ?

→ simplified Bayes Filter:

$$p_t(m) := P(m \mid x_{0:t}, y_{0:t}) \propto P(y_t \mid m, x_t) p_{t-1}(m)$$

(no transition model: assumption that map is constant)

- In the case of landmarks (FastSLAM):

$$m = (\theta_1, \dots, \theta_N)$$

$\theta_j$  = position of the  $j$ -th landmark,  $j \in \{1, \dots, N\}$

$n_t$  = which landmark we observe at time  $t$ ,  $n_t \in \{1, \dots, N\}$

We can use a separate (Gaussian) Bayes Filter for each  $\theta_j$   
(conditioned on  $x_{0:t}$ , each  $\theta_j$  is independent from each  $\theta_k$ )

$$P(\theta_{1:N} \mid x_{0:t}, y_{0:n}, n_{0:t}) = \prod_j P(\theta_j \mid x_{0:t}, y_{0:n}, n_{0:t})$$

# Particle likelihood in SLAM

- Particle likelihood for Localization Problem:

$$w_t^i = P(y_t | x_t^i)$$

(determines the new importance weight  $w_t^i$ )

- In SLAM the map is uncertain  $\rightarrow$  each particle is weighted with the *expected* likelihood:

$$w_t^i = \int P(y_t | x_t^i, m) p_{t-1}(m) dm$$

- In case of landmarks (FastSLAM):

$$w_t^i = \int P(y_t | x_t^i, \theta_{n_t}, n_t) p_{t-1}(\theta_{n_t}) d\theta_{n_t}$$

- Data association problem (actually we don't know  $n_t$ ):

For each particle separately choose  $n_t^i = \operatorname{argmax}_{n_t} w_t^i(n_t)$

# Particle-based SLAM summary

- We have a set of  $N$  particles  $\{(x^i, w^i)\}_{i=1}^N$  to represent the pose belief  $p_t(x_t)$
- For each particle we have a separate map belief  $p_t^i(m)$ ; in the case of landmarks, this factorizes in  $N$  separate 2D-Gaussians
- Iterate
  1. Resample particles to get  $N$  weight-1-particles:  $\{\hat{x}_{t-1}^i\}_{i=1}^N$
  2. Use motion model to get new “predictive” particles  $\{x_t^i\}_{i=1}^N$
  3. Update the map belief  $p_m^i(m) \propto P(y_t | m, x_t) p_{t-1}^i(m)$  for each particle
  4. Compute new importance weights  $w_t^i \propto \int P(y_t | x_t^i, m) p_{t-1}(m) dm$  using the observation model and the map belief

# Demo: Visual SLAM

- Map building from a freely moving camera



# Demo: Visual SLAM

- Map building from a freely moving camera
  - SLAM has become a big topic in the vision community..
  - features are typically landmarks  $\theta_{1:N}$  with SURF/SIFT features
  - PTAM (Parallel Tracking and Mapping) parallelizes computations...

## PTAM DTAM LSD-SLAM

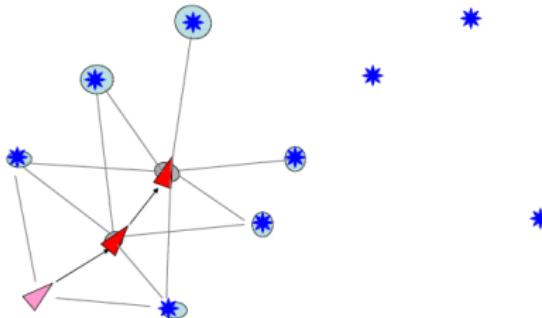
Klein & Murray: *Parallel Tracking and Mapping for Small AR Workspaces* <http://www.robots.ox.ac.uk/~gk/PTAM/>

Newcombe, Lovegrove & Davison: *DTAM: Dense Tracking and Mapping in Real-Time* ICCV 2011.

Engel, Schöps & Cremers: LSD-SLAM: Large-Scale Direct Monocular SLAM, ECCV 2014.

[http://vision.in.tum.de/\\_media/spezial/bib/engel14eccv.pdf](http://vision.in.tum.de/_media/spezial/bib/engel14eccv.pdf)

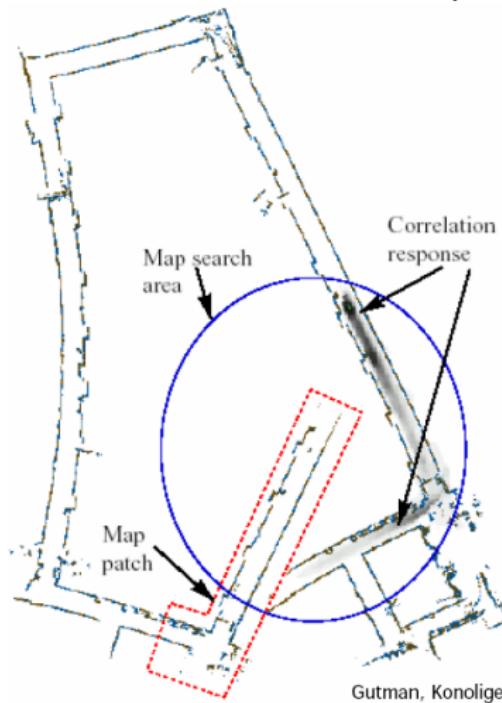
## Alternative SLAM approach: Graph-based



- Represent the previous trajectory as a graph
  - nodes = estimated positions & observations
  - edges = transition & step estimation based on scan matching
- Loop Closing: check if some nodes might coincide → new edges
- Classical Optimization:  
The whole graph defines an optimization problem: Find poses that minimize sum of edge & node errors

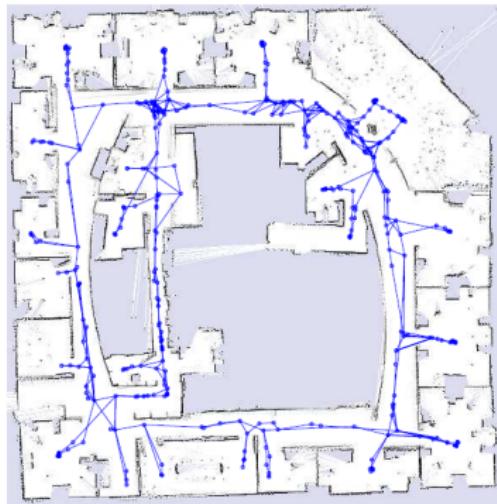
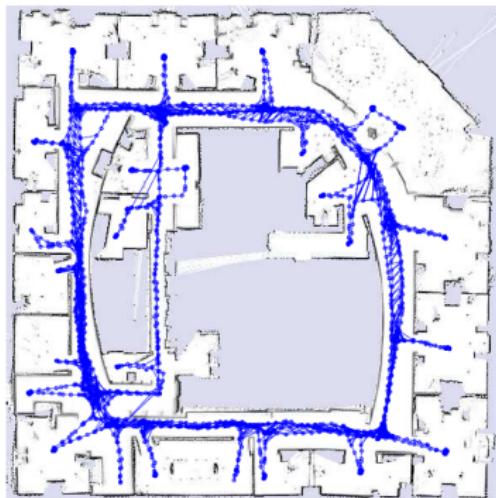
# Loop Closing Problem

(Doesn't explicitly exist in Particle Filter methods: If particles cover the belief, then “data association” solves the “loop closing problem”)



Gutman, Konolige

# Graph-based SLAM



*Life-long Map Learning for Graph-based SLAM Approaches in Static Environments*  
Kretzschmar, Grisetti, Stachniss

# Additional material

- *Robot Mapping Course*, Cyrill Stachniss, Freiburg, (slides + videos):  
<http://ais.informatik.uni-freiburg.de/teaching/ws13/mapping/>
- *Sebastian Thrun, Wolfram Burgard, Dieter Fox: Probabilistic Robotics*  
[probabilistic-robotics.org](http://probabilistic-robotics.org)

